
REINFORCEMENT LEARNING using OpenAI Gym

Uneeth Akula
50317480
Dept. of Computer Science
University at Buffalo
Buffalo, NY 14221
uneethak@buffalo.edu

Abstract

The task of this project is to build a reinforcement learning agent to navigate the classic 5x5 grid environment. I implemented reinforcement learning algorithm-Q-Learning. The agent will learn an optimal policy through Q-Learning which will allow it to take actions to reach a goal while avoiding obstacles. It finds the shortest path to reach the destination which is goal with maximum reward. The environment and agent is be built to be compatible with OpenAI Gym environments and will run effectively on computationally-limited machines.

1 Introduction

1.1 Reinforcement Learning

Reinforcement learning is an area of machine learning concerned with how software agents ought to take actions in an environment in order to maximize some notion of cumulative reward. Reinforcement learning is one of three basic machine learning paradigms, alongside supervised learning and unsupervised learning.

It differs from supervised learning in not needing labelled input/output pairs be presented, and in not needing sub-optimal actions to be explicitly corrected. Instead the focus is on finding a balance between exploration (of uncharted territory) and exploitation (of current knowledge). Reinforcement Learning lies between the spectrum of Supervised Learning and Unsupervised Learning, and there's a few important things to note:

Being greedy doesn't always work: There are things that are easy to do for instant gratification, and there's things that provide long term rewards. The goal is to not be greedy by looking for the quick immediate rewards, but instead to optimize for maximum rewards over the whole training.

Sequence matters in Reinforcement Learning: The reward agent does not just depend on the current state, but the entire history of states. Unlike supervised and unsupervised learning, time is important here.

The environment is typically stated in the form of a Markov decision process (MDP), because many reinforcement learning algorithms for this context utilize dynamic programming techniques. The main difference between the classical dynamic programming methods and reinforcement learning algorithms is that the latter do not assume knowledge of an exact mathematical model of the MDP and they target large MDPs where exact methods become infeasible.

43

44 An MDP is a 4-tuple (S, A, P, R), where

45 • S is the set of all possible states for the environment

46 • A is the set of all possible actions the agent can take

47 • $P = \Pr(s(t+1) = s' | s(t) = s, a(t) = a)$ is the state transition probability function

48 • $R: S \times A \times S \rightarrow R$ is the reward function

49 Our task is find a policy $\pi: S \rightarrow A$ which our agent will use to take actions in the environment
 50 which maximize cumulative reward, i.e.,

$$\sum_{t=0}^T \gamma^t R(s_t, a_t, s_{t+1})$$

51

52 where $\gamma \in [0, 1]$ is a discounting factor (used to give more weight to more immediate rewards), s_t
 53 is the state at time step t , a_t is the action the agent took at time step t , and s_{t+1} is the state which
 54 the environment transitioned to after the agent took the action.

55 1.2 Q-Learning:


56 Q-learning is an off-policy reinforcement learning algorithm that seeks to find the best action to
 57 take given the current state. It's considered off-policy because the q-learning function learns from
 58 actions that are outside the current policy, like taking random actions, and therefore a policy isn't
 59 needed. More specifically, q-learning seeks to learn a policy that maximizes the total reward.

60 The 'q' in q-learning stands for quality. Quality in this case represents how useful a given action is
 61 in gaining some future reward.

$$Q^{new}(s_t, a_t) \leftarrow (1 - \alpha) \cdot \underbrace{Q(s_t, a_t)}_{\text{old value}} + \underbrace{\alpha}_{\text{learning rate}} \cdot \left(\underbrace{r_t}_{\text{reward}} + \underbrace{\gamma}_{\text{discount factor}} \cdot \underbrace{\max_a Q(s_{t+1}, a)}_{\text{learned value}} \right)$$

62

63



64

65 1.3 OpenAI Gym

66 Gym is a toolkit for developing and comparing reinforcement learning algorithms. It supports
 67 teaching agents everything from walking to playing games like Pong or Pinball. The gym library
 68 is a collection of test problems — **Environments** — that you can use to work out your
 69 reinforcement learning algorithms. These environments have a shared interface, allowing you to
 70 write general algorithms. Here I am using one of the standard environment – GridEnvironment.

71 1.4 Hyper Parameters

72 **lr** or learning rate, often referred to as *alpha* or α , can simply be defined as how much you accept
 73 the new value vs the old value. Above we are taking the difference between new and old and then


74 multiplying that value by the learning rate. This value then gets added to our previous q-value
75 which essentially moves it in the direction of our latest update.

76 **Gamma** also called as discount rate. It is mainly used to calculate the future discounted reward.
77 The discount factor is mainly used to control the agent to how much area it should explore. If there
78 is no discount rate the agent simply reaches the goal but takes lot of time. If there is a discount rate
79 associated it reaches goal in quick and efficient manner by exploring all the ways. The value of
80 discount factor should be less than 1.

81 **Epsilon** is the exploration/exploitation tradeoff. It is mainly used to control the amount the
82 knowledge the agent gains. Initially we begin with larger value of epsilon by making the agent
83 learn. Then we slowly decrease the value by changing to exploitation that is the agent has enough
84 knowledge and now tries to maximize the reward.

85 **Episode** is one play that is agent moving from source to destination once. By increasing the
86 number of episodes, the agent tries to learn more and gains high rewards. So, more episodes more
87 the agent learn but at the same time with more episodes more time will be taken. Thereby
88 improving performance of the agent and gaining the reward.

89 2 Learning System



90
91 In this task we use the concept of reinforcement learning. Reinforcement Learning (RL) is one of
92 the machine Learning techniques that enables an agent to learn in an interactive environment by
93 trial and error using feedback from its own actions and experiences.

94 By following a certain sequence of steps, a learning system is selected for the problem. It allows
95 the system to automatically learn and improve from experiences. The procedure or approach
96 followed by this model is described by the steps as follows.

97 2.1 Agent

98 An agent takes actions. For eg, cab dropping off, drone making delivery, tom searching jerry. In
99 our case Agent has to choose path with maximum reward to reach the Green box.

100 2.2 Q Table:

101 When q-learning is performed we create what's called a *q-table* or matrix that follows the shape
102 of [state, action] and we initialize our values to zero. We then update and store our *q-values* after
103 an episode. This q-table becomes a reference table for our agent to select the best action based on
104 the q-value.


105 2.1 Grid Environment

106 The grid-world environment is first defined. This is the first step in implementation of the task.
107 Environment here is a 5x5 grid, so a total of 25 steps are possible. Initially the environment is run
108 just by using some random actions. This causes the agent to explore some actions thereby
109 knowing some states, actions and gains rewards. There is no particular pattern followed here.

110 There are 3 important methods in the environment-

111 `env.reset ()`: This will reset the environment and returns an initial observation.

112 `env.render ()`: For every action taken in a particular step a window will be rendered.



113

114 `env.step ()`: This returns four values that include reward, observation, a Boolean value and
115 information regarding the action taken. By using all these methods we make the agent learn from
116 the experiences and improve the performance by gaining high rewards.

117 observation (object): an environment-specific object representing your observation of the
118 environment. For example, pixel data from a camera, joint angles and joint velocities of a
119 robot, or the board state in a board game.

120 reward (float): amount of reward achieved by the previous action. The scale varies
121 between environments, but the goal is always to increase your total reward.

122 done (boolean): whether it's time to reset the environment again. Most (but not all) tasks
123 are divided up into well-defined episodes, and done being True indicates the episode has
124 terminated. (For example, perhaps the pole tipped too far, or you lost your last life.)

125 info (dict): diagnostic information useful for debugging. It can sometimes be useful for
126 learning (for example, it might contain the raw probabilities behind the environment's
127 last state change). However, official evaluations of your agent are not allowed to use this
128 for learning.

129 2.2 Making Updates

130 An agent interacts with the environment in 1 of 2 ways. The first is to use the q-table as a
131 reference and view all possible actions for a given state. The agent then selects the action based on
132 the max value of those actions. This is known as *exploiting* since we use the information we have
133 available to us to make a decision.

134 The second way to take action is to act randomly. This is called *exploring*. Instead of selecting
135 actions based on the max future reward we select an action at random. Acting randomly is
136 important because it allows the agent to explore and discover new states that otherwise may not be
137 selected during the exploitation process. You can balance exploration/exploitation using epsilon
138 (ϵ) and setting the value of how often you want to explore vs exploit.

139 The Step function just calls the Policy Function.

140 2.3 Policy function:

$$\pi(s_t) = \underset{a \in A}{\operatorname{argmax}} Q_{\theta}(s_t, a)$$

141

142 The policy function defines what step has to be taken. The epsilon is defined as 1.0 initially. The
143 function generates a random number between 0 to 1.0. If the number is less than epsilon the agent
144 uses the second method i.e., exploring. An action is selected at random. If the number is greater
145 than epsilon, the agent uses the first method i.e., exploiting. The agent selects the max value among
146 the actions from the information available to us.

147 The epsilon value is multiplied with decay, 0.996 and compared with a min constant, 0.01 in my
148 case. The maximum value is selected for each episode i.e. the epsilon is exponentially decayed and
149 the reward is calculated. When the epsilon value is high, exploration is high. As the epsilon

150 decays, the randomness decreases and the action with maximum reward is picked, i.e., exploitation
151 is high.

152 2.4 Update f unction:

153 The q_table is updated with the formula given in section (1.2)

154 2.5 Training Process:

155 I initialized decay as 0.996 and episodes as 800. We are training the model with 800 episodes. For
156 each episode the environment is reset and the actions are performed until the destination is
157 reached. The epsilon value(decayed) is appended to the epsilons list for each episode and the
158 reward is set to zero each time. Until the destination is not reached, the step function is called
159 which in-turn calls the policy function, and the action to be performed is returned. The reward
160 obtained is added to the overall reward calculated for the episode. Then I updated the q_table with
161 the previous state, next_state and the reward obtained. Then the total reward calculated is
162 appended to the rewards list. The Epsilon List and Rewards List is used for total rewards vs
163 episode graph. The epsilon is calculated by multiplying with decay and min of the calculated value
164 and 0.01 is chosen as epsilon (Exponential Decay.)

165


166 3 Results:

167 The Hyper parameters are learning rate (alpha), discount rate(gamma), epsilon and Episode. I
168 trained the model with various hyper parameters and finally got optimum path with epsilon=1.0,
169 lr=0.1, gamma=0.8, episodes=1400, decay=0.996.

170 With episodes =1400, decay=0.999 with other hyper parameters same as above

171 Did not converge


172 Epsilon Decay



173


174

Reward vs Episode



175 With episodes=1400, decay=0.991 Early Convergence

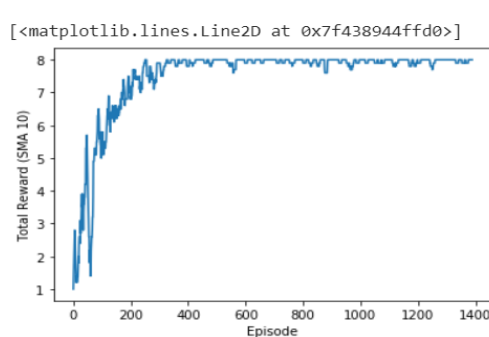
176 Epsilon Decay



177

178


Reward vs Episode



179

180 With episodes=1200 and decay=0.996 Optimum result is obtained.


181 Epsilon Decay



182

183

Reward vs Episode



184 Q_table:

```
[[ [ 5.6953279  3.96902407  5.35706371  3.97830536]
   [ 5.15015515  1.96460207  3.12193809  2.46434057]
   [ 4.5520628  0.51035565  1.01060649  0.46515947]
   [ 0.1        -0.25884837  1.05913464  0.41820297]
   [ 1.35580676 -0.40673677 -0.37418234 -0.18252754] ]

[[ [ 3.62422902  3.86461439  5.217031  3.46621215]
   [ 4.09798096  3.18028645  4.68559  3.4882741 ]
   [ 4.0951  2.74065283  2.83629613  3.11639693]
   [ 0.59901354 -0.18964892  2.88084433  0.39172489]
   [ 2.42684049 -0.38934159 -0.16693384 -0.0631441 ] ]

[[ [ 1.7607087  0.81888524  3.69795013  0.12476528]
   [ 3.90041607  1.44031874  2.15040194  0.14253487]
   [ 2.89528772  2.44343279  3.439  1.94102891]
   [ 2.71  1.02783178  2.09306921  1.83948222]
   [ 1.80947757 -0.022252  0.10274282  0.10698385] ]

[[ [ 1.6904757 -0.11004887  0.94913429 -0.36573544]
   [ 3.34997278  0.43203623  1.51095415 -0.01391113]
   [ 2.64156656  0.38319883  0.60970667  0.21217782]
   [ 1.9  1.12582835  1.57130081  0.94176705]
   [ 0.9774716 -0.08515917 -0.0631441 -0.07561 ] ]

[[ [-0.1112454 -0.074071  1.88423063 -0.29028758]
   [-0.17705596  0.30858445  2.68166717  0.01200704]
   [ 0.12286501  0.27184124  1.8992994 -0.02131716]
   [-0.1384689  0.41675677  1.  0.22066389]
   [ 0.  0.  0.  0. ] ]]
```

185

186



187

188



189

190

191 Optimal Path:





192





193

194

195



196




197

198

199

200



201

202

203 4 Conclusion:

204 Finally, the agent learns well and performs in a better way for the parameters tuned for the above
205 environment and model. By changing the value of hyper parameters we compute the reward and
206 the total time taken and then choose the values that give best reward.

207

208 5 References

209 [1] https://en.wikipedia.org/wiki/Reinforcement_learning Reinforcement Learning

210 [2] <https://towardsdatascience.com/simple-reinforcement-learning-q-learning-fcddc4b6fe56>
211 Q-Learning

212 [3] <https://gym.openai.com/docs/> OpenAI Gym

213 [4] <https://www.learndatasci.com/tutorials/reinforcement-q-learning-scratch-python-openai-gym/>
214 Reinforcement Q-Learning from Scratch in Python with OpenAI Gym

215 [5] https://www.google.com/search?q=reinforcement+learning&rlz=1C1CHBF_enUS864US864&sxsrf=ACYBGNQ3x6Y4bIJ_1UiJZ3Aw3_KaezF9-Q:1575513969203&source=lnms&tbn=isch&sa=X&ved=2ahUKEwi67-Pzvp3mAhWhrVkKHYYXUB3EQ_AUoA3oECBIQBQ&biw=2304&bih=1010#imgrc=PKBmGSfjU0AvGM: Q-Learning Image
216
217
218
219